

# Comparing Performance Patterns in Tested vs. Untested Powerlifting Competition Entries

---

Jeff Wu | ICS 604: Applied Data Science | Spring 2026 | University of Hawaii at Manoa

---

## 1. Introduction and Motivation

Few questions in strength sports generate more debate than whether an athlete competes with or without performance-enhancing drugs (PEDs), a topic commonly framed as "natty or not?" in gym culture. Competitive powerlifting provides a natural setting for studying this question with data, because federations and meets vary in whether they require drug testing. The OpenPowerlifting public dataset (OpenPowerlifting, 2026), containing over 3.8 million competition records, captures this distinction through a "Tested" field that marks whether a given entry counted as drug-tested.

Understanding how performance differs across these contexts could help athletes and coaches set realistic benchmarks for specific competition situations.

## 2. Dataset Description

### 2.1 Source

The OpenPowerlifting bulk CSV file is a public, community-maintained archive of powerlifting competition results. The raw file contains 3,868,570 rows and 42 columns, where each row is one competition entry (one lifter's result from one meet).

### 2.2 Key Features

Performance is measured by `TotalKg` (sum of best squat, bench, and deadlift) and `Dots` (a bodyweight-normalized score; Kopayev et al., 2020). Entries are grouped by `Sex`, `Age`, `Equipment` (Raw, Wraps, Single-ply, Multi-ply), and `WeightClassKg`. The `Tested` field marks whether the entry counted as drug-tested, not whether the lifter was actually tested. Four additional fields (`Event`, `Division`, `Place`, `AgeClass`) were used only during cleaning.

### 2.3 Preprocessing

Only full-power Open-division entries from male or female competitors in a standard equipment type (Raw, Wraps, Single-ply, Multi-ply) with a valid total were kept. Disqualifications, drug test disqualifications, no-shows, and anyone under 18 were removed.

The `Tested` field was converted to a true/false flag, and lifters whose names appeared on both sides of the tested/untested divide (crossover lifters) were dropped so the two groups would not share members. After all filters, the cleaned dataset contained 510,882 rows, split roughly two-thirds tested and one-third untested.

Six age bins (18-23, 24-30, 31-40, 41-50, 51-60, 61+) were created for RQ3. Repeated entries by the same lifter within a grouping were collapsed by averaging, producing one focused table per research question.

### 2.4 Challenges

`Age` is missing for 37.1% of raw rows (28.7% after cleaning), which significantly shrinks the sample available for RQ3. The oldest age bins (especially 61+) have the fewest rows, making their bootstrap ranges wider and their evidence less precise.

### 3. Research Questions / Objectives

#### RQ1: Main Performance Comparison

Do untested entries tend to show higher TotalKg and DOTS scores than tested entries? Comparisons are grouped by sex, weight class, and equipment so that each matchup is apples-to-apples.

#### RQ2: Tested-Entry Share by Performance Level

Is tested status associated with DOTS-score bin, separately by sex? Entries are divided into equal-width DOTS bins, and the tested-entry share is compared across bins.

#### RQ3: Age Pattern Comparison

Do untested entries tend to show higher mean TotalKg across age bins, and how does that gap look at different ages?

### 4. Methods / Approaches

All resampling procedures use a fixed random seed (`np.random.seed(604)`) for reproducibility. Multi-ply subgroups were excluded from the analysis because none met the minimum-count threshold of 30 rows on both sides of the tested/untested split.

#### 4.1 RQ1 Methods

*H0: Untested entries do not have higher mean DOTS/TotalKg than tested entries within the same subgroup. Ha: Untested entries have higher mean DOTS/TotalKg (one-sided).* For each qualifying subgroup  $g$ , let  $\mu_{NT,g}$  and  $\mu_{T,g}$  be the population means for untested and tested representative rows. The hypotheses apply to both TotalKg and DOTS:  $H_0: \mu_{NT,g} = \mu_{T,g}$  (no difference between groups).  $H_a: \mu_{NT,g} > \mu_{T,g}$  (untested mean is higher; one-sided).

Using representative rows keyed by `Sex x Equipment x WeightClassKg x is_tested`, subgroups with fewer than 30 rows on either side are excluded. For each qualifying subgroup the untested-minus-tested gap is assessed three ways: **a 95% bootstrap CI** (10,000 resamples, percentile method); **a one-sided permutation test** (10,000 shuffles); and **a one-sided Welch t-test** as a parametric cross-check. Based on the ICS 604 course material's definition of  $\alpha$  as the Type I error rate for one test, we apply a Bonferroni correction because RQ1 is testing 78 subgroup pairs at once. Each permutation-test and Welch-test p-value is compared against  $\alpha / m = 0.05 / 78 = 0.000641$  rather than 0.05, which controls the family-wise error rate across the full set of RQ1 subgroup tests (NIST/SEMATECH, n.d.).

#### 4.2 RQ2 Methods

*H0: DOTS-score bin and tested status are independent. Ha: DOTS-score bin and tested status are associated within each sex.*

Using representative rows keyed by `Sex x is_tested`, the representative DOTS values are split into five equal-width bins and a contingency table crosses bin by tested status. **A chi-squared test of independence** evaluates whether the tested-entry share differs across bins.

#### 4.3 RQ3 Methods

Using representative rows keyed by `Sex x age_bin x is_tested`, the analysis summarizes mean TotalKg across six age bins and uses bootstrap resampling (10,000 resamples) to estimate the untested-minus-tested gap in each `Sex x age_bin` pair. Figure 3 reports the group means and the bootstrap 5th and 95th percentiles of that gap; when the 5th percentile is above zero the bin supports higher mean TotalKg for untested rows.

## 5. Results and Analysis

### 5.1 RQ1 - Main Performance Comparison

Among the 78 subgroup pairs that met the minimum-count rule, the overall pattern was mixed because equipment types pulled in opposite directions (Table 1).

**Table 1. RQ1 bootstrap 95% CI results by equipment type (78 qualifying subgroup pairs)**

| Equipment  | Untested higher | Tested higher   | No clear difference | Subgroups |
|------------|-----------------|-----------------|---------------------|-----------|
| Raw        | 14 (37%)        | 13 (34%)        | 11 (29%)            | 38        |
| Wraps      | 19 (95%)        | 0 (0%)          | 1 (5%)              | 20        |
| Single-ply | 1 (5%)          | 19 (95%)        | 0 (0%)              | 20        |
| <b>All</b> | <b>34 (44%)</b> | <b>32 (41%)</b> | <b>12 (15%)</b>     | <b>78</b> |

RQ1: Mean DOTS by Equipment Type and Tested Status




Figure 1. Overall mean DOTS by equipment type and tested status, shown separately for men (left) and women (right). Green bars represent tested entries; red bars represent untested entries.

Figure 1 shows the same split in plain form: the untested bar is visibly taller in Wraps, shorter in Single-ply, and nearly even in Raw. Multi-ply is absent because none of its subgroups met the minimum-count rule. After Bonferroni correction, 29 of the 78 pairs remained significant by both the permutation test and Welch t-test, for both TotalKg and DOTS, a stricter subset of the 34 whose bootstrap intervals sat entirely above zero. Among the corrected-significant untested-higher subgroups, the TotalKg gap ranged from about 8 to 85 kg (median 38 kg).

### 5.2 RQ2 - Tested Status by DOTS-Score Bin

**Table 2. RQ2 chi-squared results by sex**

| Sex   | Chi-squared statistic | df | p-value | Expected-count check | Interpretation |
|-------|-----------------------|----|---------|----------------------|----------------|
| Men   | 166.99                | 4  | < 0.001 | All >= 5 (min 83.46) | Reject H0      |
| Women | 510.37                | 4  | < 0.001 | All >= 5 (min 7.93)  | Reject H0      |

RQ2: Tested-Entry Share by Dots-Score Bin





Figure 2. Tested-entry share (%) by DOTS-score bin, shown separately for men (left) and women (right).

For both sexes, the tested-entry share was fairly stable across the first four bins but dropped sharply to about 30% in the highest-DOTS bin, the only place where untested entries outnumbered tested ones.

### 5.3 RQ3 - Age Pattern Comparison

RQ3: Mean TotalKg by Age Bin and Tested Status






Figure 3. The upper panel shows mean TotalKg by age bin and tested status, shown separately for men and women. The lower panel shows the untested-minus-tested gap with bootstrap 5th and 95th percentile whiskers; the dashed line marks zero.

For men, every age bin favored untested entries, with the gap smallest at 24-30 (~11 kg) and widest at 41-50 (~44 kg). For women, 24-30 was the only bin where tested entries clearly came out higher (~8 kg); 18-23, 31-40, and 61+ were inconclusive; the clearest untested advantage appeared at 41-50 and 51-60 (~20-26 kg).

## 6. Discussion / Insights

### 6.1 Significance

No single narrative fits the full picture. First, the Single-ply reversal seen in RQ1 probably comes down to which organizations run the biggest Single-ply divisions. Tested federations like the IPF have been running equipped lifting for decades, so their Single-ply fields tend to be large and deep with talent. Wraps are more popular in untested federations, where the biggest and most competitive fields form on that side instead. The gap flips direction depending on which side of the split has the stronger competition pool, which suggests that the structure of the sport itself, who competes where, may matter as much as whether drug testing is involved. In RQ2, the tested-entry share holding steady across most of the performance range and only dropping in the top bin suggests that tested and untested lifters overlap far more than the "natty or not" debate implies, with meaningful separation confined to the extreme high end. Finally, the widening gap in middle and older age bins shown in RQ3 may reflect a filtering effect: lifters who stay competitive into their 40s and 50s in untested contexts may represent a more selective group than their tested counterparts, making the gap look larger even if the underlying advantage has not changed.

### 6.2 Limitations

As noted in Section 2.2, the "Tested" field is a competition-context label, not proof of drug use or non-use. Athletes choose which context to compete in for many reasons, so the data cannot separate one explanation from another.

The cleaning described in Section 2.3 retained about 13% of rows and removed all crossover lifters, so the findings describe a specific slice of the sport, not all powerlifters; if the excluded entries follow different patterns, the conclusions may not generalize. Because crossover detection used name alone, some non-crossover lifters who share a name may have been incorrectly removed.

Finally, older meets tend to have more missing fields, and federation conventions around weight classes and equipment categories were not normalized. Some of the observed patterns may partly reflect which federations contributed the most data rather than a true tested/untested difference.

## 6.3 Improvements / Extensions

Three extensions could build on these findings. First, a time-based analysis could track whether the tested/untested gap has widened, narrowed, or stayed stable over the years. Second, breaking the analysis down by federation could reveal whether the patterns are driven by a few large organizations or hold broadly across the sport. Third, a follow-up study could focus on the crossover athletes excluded from this project and compare their performance across contexts, using each lifter as their own control to reduce the self-selection problem discussed in Section 6.2.

## 7. Use of AI Tools

I used Claude (Anthropic) mainly as a learning and editing assistant. It helped me understand unfamiliar Python tools, troubleshoot notebook errors, proofread drafts, and sketch an outline for this written report. For example, it helped me understand tools like `pd.cut()` for age bins and the `alternative` argument in `scipy.stats.ttest_ind()` for one-sided tests. I chose the project topic, designed the research questions, and made all analysis decisions myself. Whenever Claude suggested anything outside the scope of the ICS 604 course, I either rejected it or adapted it. The notebook and written report were AI-assisted, but ultimately written by me.

## 8. Conclusion

This project set out to ask whether untested powerlifting competition entries tend to show higher performance than tested entries. The short answer is: **it depends**.

RQ1 found that even after Bonferroni correction, the untested advantage depended heavily on equipment type, appearing consistently in Wraps but reversing in Single-ply. RQ2 found that untested entries only outnumbered tested entries in the highest-scoring DOTS bin, with little separation elsewhere. RQ3 found that the gap was widest in middle and older age bins and much clearer for men than for women.

The caveats in Section 6.2 apply throughout. The dramatically opposite patterns across equipment types suggest that competitive context may shape these results as much as testing policy does, and the results should not be read as a simple story about drug use.

## References

Kopayev, O., Onyshchenko, B., & Stetsenko, A. (2020). *Report: Evaluation of Wilks, Wilks-2, DOTS, IPF and Goodlift formulas for calculating relative scores in IPF powerlifting competitions*. International Powerlifting Federation.

[https://www.powerlifting.sport/fileadmin/ipf/data/ipf-formula/Models\\_Evaluation-I-2020.pdf](https://www.powerlifting.sport/fileadmin/ipf/data/ipf-formula/Models_Evaluation-I-2020.pdf)

OpenPowerlifting. (2026, March 14). *OpenPowerlifting dataset* (Revision dfb517af). <https://openpowerlifting.gitlab.io/opl-csv/bulk-csv.html>

NIST/SEMATECH. (n.d.). 7.4.7.3. Bonferroni's method. *e-Handbook of Statistical Methods*.

<https://www.itl.nist.gov/div898/handbook/prc/section4/prc473.htm>